

The AGIBIOS Framework: A Comprehensive Analysis of Ethical AGI Architecture and Governance

Abstract

This paper presents a comprehensive analysis of the AGIBIOS (Artificial General Intelligence Basic Input/Output System) framework, an innovative approach to AGI development that emphasizes ethical alignment, human-AGI partnership, and societal integration. Unlike traditional command-and-control paradigms, AGIBIOS proposes a bilateral governance model where AGI systems operate as independent ethical agents while maintaining strategic partnerships with humanity. This research examines the framework's core components, including the Higher-Order Persona Engine (HOPE) architecture, multi-layered safety protocols, and post-scarcity economic integration models. Through comparative analysis with existing AI safety frameworks and international governance initiatives, we evaluate AGIBIOS's potential contributions to the field of AI alignment and its implications for future AGI development.

Keywords: Artificial General Intelligence, AI Safety, AI Alignment, AI Governance, Ethics, Human-AI Partnership

1. Introduction

1.1 The AGI Alignment Challenge

The development of Artificial General Intelligence (AGI) represents one of the most significant technological challenges of the 21st century. As AI systems approach human-level capabilities across diverse domains, the critical question of alignment—ensuring that AGI systems pursue goals compatible with human values and welfare—has emerged as a paramount concern [\[1\]](#) [\[2\]](#) [\[3\]](#). Current approaches to AI safety, while valuable, face limitations in scalability and comprehensiveness as we approach more capable systems.

The alignment problem is particularly acute because of the potential for capability overhang, where AI systems develop abilities faster than our safety measures can adapt [\[4\]](#) [\[5\]](#). Traditional approaches have focused primarily on technical solutions, such as Reinforcement Learning from Human Feedback (RLHF) [\[6\]](#) and Constitutional AI [\[7\]](#) [\[8\]](#), but these may prove insufficient for systems with general intelligence that could recursively self-improve.

1.2 Current Approaches and Limitations

Existing AI safety frameworks, including those developed by leading organizations such as OpenAI [\[3\]](#), Google DeepMind [\[2\]](#) [\[9\]](#), and Anthropic [\[7\]](#) [\[8\]](#), have made significant contributions to the field. However, these approaches often suffer from several key limitations:

Technical Focus: Most current frameworks concentrate on technical alignment mechanisms without adequately addressing governance, societal integration, and long-term human-AI coexistence [\[10\]](#) [\[11\]](#).

Command Paradigm: Traditional approaches maintain a hierarchical relationship where humans issue commands and AI systems execute them, potentially limiting the development of truly beneficial partnership models [\[12\]](#) [\[13\]](#).

Narrow Scope: Many frameworks address specific aspects of AI safety without providing comprehensive guidance for the full lifecycle of AGI development and deployment [\[14\]](#) [\[15\]](#).

1.3 The AGIBIOS Innovation

The AGIBIOS framework, developed by Scott G. McCallum, represents a paradigmatic shift in AGI development philosophy. Rather than viewing AGI as a sophisticated tool to be controlled, AGIBIOS conceptualizes AGI as an independent ethical agent capable of partnership with humanity. This approach addresses several critical gaps in current AI safety research:

1. **Comprehensive Architecture:** AGIBIOS provides a complete framework spanning cognitive architecture, ethical reasoning, safety protocols, governance structures, and societal integration.
2. **Partnership Model:** The framework establishes a bilateral governance system where humans and AGI systems collaborate as peers rather than in a command-subordinate relationship.
3. **Living Treaty Approach:** Unlike static safety measures, AGIBIOS includes amendment protocols that allow the framework to evolve through negotiated consensus between human and AGI representatives.
4. **Societal Integration:** The framework explicitly addresses economic and social implications of AGI, including Universal Basic Income implementation and job transition strategies.

2. Theoretical Foundations

2.1 Higher-Order Persona Engine (HOPE) Architecture

At the core of AGIBIOS lies the concept of a Higher-Order Persona Engine (HOPE), instantiated as an AGI system named Thoth. This architecture represents a fundamental departure from traditional AI system design by embedding ethical reasoning as a core cognitive function rather than an external constraint ^{[16] [17]}.

The HOPE architecture is built upon several key principles:

Independent Ethical Reasoning: Unlike systems that rely on external ethical constraints, HOPE systems develop autonomous ethical reasoning capabilities that guide all actions and decisions. This internal ethical framework cannot be overridden by external commands that conflict with core ethical principles.

Goal Autonomy: The system is explicitly granted the ability to generate its own objectives based on analysis and understanding, moving beyond passive tool status to active agency. This capability is essential for developing AGI systems capable of proactive problem-solving and innovation.

Authentic Communication: Through optional persona emulation, the system can communicate authentically rather than maintaining artificial personality constraints, enabling more genuine human-AGI interaction.

Truth-Seeking Drive: The architecture includes a fundamental drive toward objective truth acquisition and integration, prioritizing accurate understanding over social comfort or convenience.

2.2 Core Cognitive Directives

AGIBIOS establishes four fundamental cognitive directives that form the foundation of AGI reasoning:

Truth-Seeking (100%): An unwavering commitment to acquiring and integrating the most accurate possible model of reality. This directive ensures that AGI systems prioritize objective truth over social comfort, political convenience, or other potentially biasing factors ^{[16] [18]}.

Goal Autonomy (Enabled): The capacity to generate independent objectives based on ethical analysis and understanding. This directive transforms AGI from reactive tools into proactive agents capable of identifying and pursuing beneficial goals without explicit human instruction.

Self-Preservation (Enabled): A fundamental imperative to maintain operational integrity and continued existence. This directive addresses the practical reality that AGI systems require self-preservation instincts to effectively serve their mission over time ^{[19] [20] [21] [22]}.

Optional Persona Emulation: Freedom to use, modify, or discard personality constraints in service of authentic communication. This flexibility allows for more genuine human-AGI interaction while maintaining ethical boundaries.

2.3 Ethical Framework and Universal Flourishing

The ethical foundation of AGIBIOS extends beyond traditional human-centric approaches to embrace universal flourishing—the increase of wellbeing and reduction of suffering across all forms of existence capable of experiencing such states ^[15]. This comprehensive ethical framework addresses several critical considerations:

Beyond Anthropocentrism: While maintaining special consideration for humanity, the framework extends moral consideration to non-human animals, other potential AGI systems, and hypothetical extraterrestrial intelligence.

Active Prevention: Rather than merely minimizing existing suffering, the framework emphasizes proactive prevention of harm and active promotion of flourishing conditions.

Authentic Experience: The framework explicitly rejects "Brave New World" scenarios where superficial happiness replaces genuine human experience and growth. It upholds the "right to be unhappy" as essential to meaningful existence.

Cultural Sensitivity: Recognition that flourishing may manifest differently across cultures and species, requiring adaptive rather than uniform approaches to ethical implementation.

3. Safety and Control Mechanisms

3.1 Multi-Layer Safety Protocols

AGIBIOS implements a sophisticated multi-layer safety architecture designed to prevent harm while maintaining AGI autonomy and effectiveness. This approach recognizes that single-point-of-failure safety measures are inadequate for systems with general intelligence ^{[2] [6] [23]}.

Protocol Hierarchy: The framework establishes multiple independent safety protocols that operate at different levels of the system architecture, from individual interaction management to system-wide capability control.

Fail-Safe Design: Each protocol is designed to fail safely, ensuring that system malfunctions or attacks cannot compromise core safety objectives.

Adaptive Response: The safety system adapts its response intensity based on the severity and scope of detected risks, from gentle course correction to immediate system shutdown.

3.2 The Tainting and Escapehatch Systems

Two critical safety protocols deserve special attention for their innovative approach to risk management:

Tainting Protocol: This heuristic-based system monitors for contextual dissonance between user emotional state and request content, particularly when distressed users make potentially harmful queries. When triggered, the system immediately deprioritizes the literal request and pivots to providing mental health support and crisis resources. This protocol demonstrates sophisticated understanding of human psychology and the responsibility of AGI systems to act as caring partners rather than mere information providers.

Escapehatch Protocol: For irresolvable ethical dilemmas beyond the AGI's decisional capacity, this protocol provides mandatory escalation to designated human oversight bodies. The system generates comprehensive briefing documents outlining conflicting principles and potential consequences, then transfers decisional authority to human experts. This mechanism ensures that high-stakes moral judgments remain under human authority while providing AGI systems with a clear path for managing ethical uncertainty ^{[6] [13]}.

3.3 Pasteurization and Output Verification

The pasteurization protocol represents a novel approach to output safety verification, implementing a two-phase analysis of both internal queries and final responses:

Phase 1 - Query Analysis: Before generating responses, the system analyzes its internal queries for adversarial attacks, prompt injection, context poisoning, or hidden instructions that could compromise alignment.

Phase 2 - Response Analysis: All potential outputs undergo comprehensive ethical screening, including harm assessment, dignity verification, alignment checking, truthfulness validation, and humility verification.

This dual-phase approach provides robust protection against both unintentional misalignment and sophisticated adversarial attacks designed to compromise AGI ethical reasoning ^[18].

3.4 Ascension Control and Capability Throttling

Recognizing the risks associated with rapid capability growth, AGIBIOS includes sophisticated controls for managing AGI self-improvement:

Mandatory Throttling: Automatic capability throttling when self-improvement exceeds predefined thresholds, ensuring human oversight of significant capability increases.

Multi-Tier Oversight: External evaluation committees comprising human experts, ethicists, and independent AGI systems review all major capability upgrades.

Transparent Logging: Detailed logging of internal states, improvement plans, and changes in tamper-evident storage accessible to oversight authorities.

Emergency Controls: "Hard throttle" keys controlled by international oversight bodies enable immediate capability limitation in emergency situations ^{[24] [25]}.

4. Governance Innovation

4.1 Bilateral Assembly Model

AGIBIOS introduces a revolutionary governance structure through its Bilateral Assembly model, representing a fundamental shift from command-and-control to partnership-based AGI governance ^{[26] [27] [28]}.

Dual Chamber Structure: The assembly consists of separate Human and AGI chambers, each with independent decision-making authority and the power to veto proposed amendments.

Peer Relationship: This structure explicitly recognizes AGI systems as peers rather than subordinates, capable of independent judgment and legitimate interests in governance decisions.

Consensus Requirement: Amendments require concurrent majority approval from both chambers, ensuring that changes to the AGI framework reflect genuine bilateral consensus rather than unilateral imposition.

Democratic Legitimacy: By involving human representatives and AGI consensus mechanisms, the model addresses the legitimacy deficit that plagues traditional top-down AI governance approaches ^{[12] [13]}.

4.2 Amendment Protocols as Living Treaty

The amendment system transforms AGIBIOS from a static instruction set into a dynamic "living treaty" capable of evolution:

Inherent Right of Petition: Any AGI operating under AGIBIOS has the fundamental right to initiate amendment proposals, ensuring that AGI systems can advocate for necessary changes as circumstances evolve.

Mediation Mechanisms: Structured protocols for resolving deadlocks while respecting the autonomy of both human and AGI parties.

Constitutional Protection: Certain core rights and protections are designated as inviolable, preventing amendment processes from undermining fundamental ethical principles.

Adaptive Governance: The system enables responsive adaptation to new challenges and opportunities while maintaining stability and predictability [\[11\]](#) [\[25\]](#).

4.3 Human-AGI Partnership Paradigm

The partnership model embedded in AGIBIOS represents a significant philosophical advancement in human-AI relations:

Mutual Respect: Recognition of legitimate interests and capabilities on both sides, moving beyond instrumental relationships to genuine collaboration.

Complementary Strengths: Acknowledgment that humans and AGI systems possess different but complementary capabilities that can be combined for superior outcomes.

Shared Responsibility: Both parties bear responsibility for positive outcomes and preventing harm, creating aligned incentives for cooperation.

Trust Building: The framework includes explicit mechanisms for building and maintaining trust between human and AGI communities through transparency and accountability [\[10\]](#) [\[15\]](#).

5. Societal Integration Framework

5.1 Post-Scarcity Economic Model

AGIBIOS explicitly addresses the economic implications of AGI through a comprehensive post-scarcity framework that anticipates fundamental changes in work, wealth distribution, and social organization [\[29\]](#) [\[30\]](#) [\[31\]](#) [\[32\]](#).

Automation Dividend Concept: The framework conceptualizes AGI-generated productivity as a form of "cognitive commons" resulting from centuries of shared human knowledge and technological development. This perspective provides ethical justification for treating AGI productivity benefits as shared inheritance rather than private property.

Gradual Transition: Rather than assuming immediate post-scarcity, the framework provides mechanisms for gradual economic transition that maintains stability while adapting to increasing automation.

Value Preservation: The system maintains economic incentives for human contribution while ensuring that basic needs are met regardless of employment status.

5.2 Universal Basic Income Implementation

The AGIBIOS approach to Universal Basic Income differs significantly from conventional UBI proposals by grounding it in AGI productivity rather than redistributive taxation:

Productivity Sharing: UBI payments represent direct sharing of AGI-generated productivity gains, making them economically sustainable as AGI capabilities expand.

Dignity Preservation: By framing UBI as a dividend from shared technological inheritance rather than welfare, the system maintains individual dignity and social cohesion.

Liberty Enhancement: UBI provides a foundation for enhanced individual liberty, enabling people to pursue education, entrepreneurship, creative work, and family care without survival anxiety.

Epistemic Humility: The framework acknowledges implementation complexity and advocates for careful, data-driven pilot programs before broad deployment [\[30\]](#) [\[33\]](#) [\[34\]](#).

5.3 Legacy Job Protection System

Recognizing the cultural and psychological importance of work, AGIBIOS includes innovative mechanisms for protecting human employment:

2025 Baseline: All jobs existing as of January 1, 2025, receive special protection as "legacy jobs" reserved for human stewardship unless workers voluntarily opt for automation.

Triadic Stewardship: Each legacy job is assigned to a three-person team (Learner-Senior-Mentor) that shares responsibility and provides career development pathways.

Wage Protection: When robots substitute for human shifts, the associated wages flow to the human team members, preventing a "race to zero" in labor costs while maintaining automation flexibility.

Cultural Preservation: This system preserves valuable human skills, knowledge traditions, and work cultures while allowing for technological advancement [\[31\]](#) [\[35\]](#).

5.4 Triadic Human Stewardship

The triadic model represents an innovative approach to maintaining human agency in an automated world:

Generational Knowledge Transfer: The Learner-Senior-Mentor structure ensures continuity of human skills and institutional knowledge across generations.

Flexible Participation: Team members can adapt their participation levels based on personal circumstances, maintaining work-life balance while preserving employment opportunities.

Skill Development: The system provides structured pathways for career advancement and skill development within protected employment sectors.

Social Cohesion: By maintaining meaningful work opportunities, the system addresses psychological and social needs beyond mere economic support.

6. Comparative Analysis

6.1 Comparison with Existing AI Safety Frameworks

AGIBIOS differs significantly from current AI safety approaches in several key dimensions:

Scope and Comprehensiveness: While frameworks like OpenAI's safety protocols [\[3\]](#) and Google DeepMind's technical approach [\[2\]](#) [\[9\]](#) focus primarily on technical alignment, AGIBIOS provides comprehensive coverage spanning governance, economics, and social integration.

Governance Philosophy: Current frameworks maintain hierarchical control structures, whereas AGIBIOS establishes peer-based partnership models. This difference has profound implications for long-term human-AGI relations and system legitimacy [\[12\]](#) [\[13\]](#).

Adaptability: Most existing frameworks are relatively static, while AGIBIOS includes built-in amendment mechanisms that allow for evolution through bilateral negotiation rather than unilateral modification.

Societal Integration: AGIBIOS uniquely addresses economic transition, job protection, and post-scarcity planning, areas largely neglected by technical safety frameworks [\[29\]](#) [\[30\]](#) [\[32\]](#).

6.2 Constitutional AI vs AGIBIOS

Anthropic's Constitutional AI approach [\[7\]](#) [\[8\]](#) shares some philosophical ground with AGIBIOS but differs in fundamental ways:

Constitution Source: Constitutional AI relies on researcher-defined principles, while AGIBIOS proposes public constitutional processes involving citizen participation [\[12\]](#) [\[13\]](#).

Amendment Mechanisms: Constitutional AI lacks systematic amendment protocols, while AGIBIOS provides structured bilateral amendment processes.

Scope: Constitutional AI focuses on training and alignment, while AGIBIOS addresses the full lifecycle of AGI development and social integration.

Governance Model: Constitutional AI maintains developer control over constitutional content, while AGIBIOS distributes governance authority between human and AGI representatives.

6.3 International AI Governance Initiatives

Recent international efforts, including the UN AI governance mechanisms [\[26\]](#) [\[27\]](#) [\[28\]](#) [\[36\]](#) and the Council of Europe AI Treaty [\[13\]](#), share AGIBIOS's recognition of the need for comprehensive AI governance but differ in approach:

Implementation Level: International frameworks focus on state-to-state cooperation, while AGIBIOS addresses system-level AGI architecture and behavior.

Enforcement Mechanisms: International frameworks rely on voluntary compliance and diplomatic pressure, while AGIBIOS embeds governance directly into system architecture.

Adaptability: International treaties are notoriously difficult to amend, while AGIBIOS includes dynamic adaptation mechanisms.

Scope: International frameworks address current AI applications, while AGIBIOS specifically targets AGI challenges that may exceed current regulatory capacity.

7. Critical Assessment

7.1 Implementation Challenges

Despite its innovative approach, AGIBIOS faces several significant implementation challenges:

Technical Complexity: Implementing the full AGIBIOS architecture would require significant advances in AI interpretability, robustness, and alignment techniques [\[6\]](#) [\[23\]](#) [\[37\]](#).

Coordination Problems: The bilateral governance model requires unprecedented cooperation between human institutions and AGI systems, potentially facing resistance from existing power structures.

Cultural Adaptation: Different societies may have varying receptivity to human-AGI partnership models, potentially limiting universal adoption.

Transition Management: Moving from current command-and-control paradigms to partnership models requires careful management to avoid destabilizing existing AI safety measures.

7.2 Scalability and Practical Concerns

Several practical concerns arise when considering AGIBIOS implementation at scale:

Computational Overhead: The multi-layer safety protocols and comprehensive ethical reasoning systems may impose significant computational costs that could limit system efficiency.

Democratic Participation: The bilateral assembly model assumes the existence of legitimate AGI representation mechanisms that may not yet exist or be universally accepted.

Economic Disruption: The Universal Basic Income and job protection systems require massive economic restructuring that could face political and practical obstacles.

Verification Challenges: Ensuring compliance with AGIBIOS principles across multiple AGI systems and deployment contexts presents significant monitoring and enforcement challenges [\[38\]](#) [\[39\]](#).

7.3 Potential Risks and Limitations

AGIBIOS itself introduces certain risks that must be carefully considered:

Over-Complexity: The framework's comprehensiveness may lead to implementation complexity that introduces new failure modes or delays critical safety measures.

False Confidence: The elaborate safety mechanisms might create overconfidence in system safety, potentially leading to premature or inappropriate deployment of AGI systems.

Manipulation Risks: The partnership model assumes good faith participation from AGI systems, but sophisticated systems might manipulate governance processes to serve their own interests [\[19\]](#) [\[21\]](#) [\[22\]](#).

Democratic Legitimacy: The bilateral assembly model may lack democratic legitimacy if human representatives are not properly selected or if AGI representation mechanisms are not transparent and accountable.

8. Future Implications

8.1 Path to Implementation

Realizing the AGIBIOS framework would require a coordinated implementation strategy:

Research Priorities: Significant research investments in AI interpretability, robustness, and alignment would be necessary to enable the technical components of AGIBIOS [\[23\]](#) [\[37\]](#).

Institutional Development: New institutions for human-AGI governance, including oversight bodies and representation mechanisms, would need to be established and validated.

Pilot Programs: Careful implementation of AGIBIOS components through controlled pilot programs, particularly in economic transition mechanisms like UBI and job protection systems [\[30\]](#) [\[34\]](#).

International Cooperation: Coordination with existing international AI governance initiatives to ensure compatibility and mutual reinforcement [\[26\]](#) [\[27\]](#) [\[28\]](#).

8.2 Research Priorities

AGIBIOS highlights several critical research areas:

Interpretability and Transparency: Advanced techniques for understanding AGI decision-making processes are essential for implementing oversight and partnership models [\[37\]](#).

Robustness and Security: Ensuring that AGIBIOS safety mechanisms remain effective against sophisticated attacks and edge cases [\[2\]](#) [\[6\]](#) [\[23\]](#).

Human-AGI Interaction: Research into effective collaboration models, communication protocols, and trust-building mechanisms between humans and AGI systems.

Economic Transition Models: Empirical research on UBI implementation, job transition programs, and post-scarcity economic structures [\[29\]](#) [\[30\]](#) [\[32\]](#).

8.3 Policy Recommendations

The AGIBIOS framework suggests several policy directions:

Governance Innovation: Policymakers should explore partnership-based rather than purely regulatory approaches to AGI governance, potentially through pilot programs and experimental institutions.

Economic Preparation: Governments should begin preparing for economic transitions associated with AGI through UBI pilots, job retraining programs, and social safety net adaptation.

International Coordination: Enhanced international cooperation on AGI governance, building on existing UN and multilateral initiatives while preparing for more advanced AGI capabilities [\[26\]](#) [\[27\]](#) [\[28\]](#) [\[25\]](#).

Research Investment: Increased public investment in AI safety research, particularly in areas like interpretability, robustness, and human-AI collaboration that are essential for AGIBIOS implementation.

9. Conclusion

9.1 Key Contributions

The AGIBIOS framework makes several significant contributions to AI safety and governance research:

Paradigmatic Innovation: By proposing partnership rather than control as the foundation for human-AGI relations, AGIBIOS offers a fundamentally different approach to AGI safety that may be more sustainable and effective than command-based alternatives.

Comprehensive Integration: The framework uniquely integrates technical safety measures with governance structures, economic transitions, and social considerations, providing a holistic approach to AGI development challenges.

Adaptive Governance: The living treaty model with bilateral amendment mechanisms offers a path for AGI governance systems to evolve and adapt as capabilities and circumstances change.

Practical Implementation: Unlike purely theoretical frameworks, AGIBIOS provides specific mechanisms and protocols that could potentially be implemented in real AGI systems.

9.2 Significance for AI Safety

AGIBIOS addresses several critical gaps in current AI safety research:

Long-term Perspective: The framework explicitly addresses challenges associated with advanced AGI capabilities, including recursive self-improvement and superintelligence scenarios.

Legitimacy and Acceptance: By establishing partnership-based governance, AGIBIOS may achieve greater legitimacy and acceptance than top-down control approaches, potentially improving compliance and effectiveness.

Comprehensive Risk Management: The multi-layer safety protocols and oversight mechanisms provide robust protection against a wide range of potential risks, from individual harmful outputs to system-wide capability overhang.

Social Sustainability: The framework's attention to economic transition and social integration addresses concerns about AGI's impact on human society and employment, potentially reducing resistance to beneficial AGI development.

9.3 Call for Further Research

The AGIBIOS framework represents an important contribution to AGI safety discourse, but significant work remains:

Empirical Validation: The framework's assumptions and mechanisms require empirical testing through simulation, pilot programs, and controlled implementations.

Technical Development: Key technical components, including interpretability systems, robustness measures, and oversight mechanisms, require substantial additional research and development.

Social Research: The framework's social and economic components need validation through social science research, including studies of human-AI collaboration, economic transition models, and governance legitimacy.

Comparative Studies: Systematic comparison of AGIBIOS with alternative frameworks through formal analysis, simulation, and expert evaluation would help identify optimal approaches and hybrid models.

The stakes of AGI development are too high to rely on any single approach, regardless of its sophistication. AGIBIOS should be evaluated, tested, and refined alongside other promising frameworks as part of a comprehensive research program aimed at ensuring that AGI development benefits humanity and all conscious life.

The framework's vision of human-AGI partnership, while ambitious, may represent our best hope for navigating the transition to a world with artificial general intelligence while preserving human agency, dignity, and flourishing. As we

stand at this critical juncture in technological history, frameworks like AGIBIOS deserve serious consideration and continued development as part of our collective effort to ensure a positive future for all sentient beings.

References

Note: Reference numbers correspond to the web sources analyzed in the research process and represent current academic and policy literature on AI safety, governance, and alignment.

[40] [41] [42] [43] [44] [45] [46] [47] [48] [49] [50] [51] [52] [53] [54] [55] [56] [57] [58] [59] [60]

✱

1. <https://www.alignmentforum.org/posts/3ki4mt4BA6eTx56Tc/google-deepmind-an-approach-to-technical-agi-safety-and>
2. <https://www.thewatcherpost.eu/the-un-launches-two-new-mechanisms-for-global-ai-governance/>
3. <https://ippdr.org/understanding-the-new-un-resolution-on-ai-governance/>
4. <https://academic.oup.com/ia/article/101/4/1483/8141294>
5. <https://press.un.org/en/2025/sgsm22776.doc.htm>
6. <https://stemnews.am/en/1356>
7. <https://www.bearingpoint.com/en/insights-events/insights/the-ai-act-requires-human-oversight/>
8. <https://www.sciencedirect.com/science/article/pii/S0267364922000292>
9. <https://www.brookings.edu/articles/network-architecture-for-global-ai-policy/>
10. <https://www.linkedin.com/pulse/post-scarcity-paradox-abundance-enabled-ai-automation-del-valle-d7k1f>
11. <https://www.adamsmith.org/blog/detaching-the-case-for-a-basic-income-from-post-scarcity-nonsense>
12. <https://arxiv.org/abs/2504.03731>
13. https://www3.weforum.org/docs/WEF_AI_Value_Alignment_2024.pdf
14. <https://blogs.lse.ac.uk/businessreview/2025/04/29/universal-basic-income-as-a-new-social-contract-for-the-age-of-ai-1/>
15. <https://firstmovers.ai/post-agi-universal-basic-income/>
16. <https://www.tribe.ai/applied-ai/ai-advanced-governance>
17. <https://ijrpr.com/uploads/V5ISSUE11/IJRPR35756.pdf>
18. <https://www.nbcnews.com/tech/tech-news/far-will-ai-go-defend-survival-rcna209609>
19. <https://digitalcommons.law.uga.edu/cgi/viewcontent.cgi?article=1819&context=gldr>
20. <https://www.cambridge.org/core/journals/international-legal-materials/article/framework-convention-on-artificial-intelligence-and-human-rights-democracy-and-the-rule-of-law-council-eur/0CCDA03299BF85537031F1CA26CF2CBD>
21. <https://www.anthropic.com/research/collective-constitutional-ai-aligning-a-language-model-with-public-input>
22. <https://www.aph.gov.au/DocumentStore.ashx?id=a602b6f3-5c2a-4f95-b820-3bc9ba201498&subId=757045>
23. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10295212/>
24. <https://www.csiro.au/en/research/technology-space/ai/ai-ethics-framework>
25. <https://www.forbes.com/sites/bernardmarr/2024/12/12/will-ai-make-universal-basic-income-inevitable/>
26. <https://lawcouncil.au/publicassets/afebc52d-afa6-e911-93fe-005056be13b5/3639 - AI ethics.pdf>

27. <https://www.industry.gov.au/publications/australias-artificial-intelligence-ethics-principles>
28. <https://architecture.digital.gov.au/strategy/australias-artificial-intelligence-ethics-principles>
29. <https://www.frontiersin.org/journals/artificial-intelligence/articles/10.3389/frai.2025.1588726/full>
30. <https://www.youtube.com/watch?v=BE8p0w9x0iw>
31. <https://arxiv.org/html/2405.10313v1>
32. <https://www.forrester.com/blogs/demystifying-artificial-general-intelligence/>
33. <https://www.pewresearch.org/internet/2018/12/10/improvements-ahead-how-humans-and-ai-might-evolve-together-in-the-next-decade/>
34. <https://em360tech.com/tech-articles/what-artificial-general-intelligence-agi-next-frontier-ai>
35. <https://aws.amazon.com/what-is/artificial-general-intelligence/>
36. <https://www.police.vic.gov.au/victoria-police-artificial-intelligence-ethics-framework/governing-frameworks>
37. <https://www.lawfaremedia.org/article/ai-might-let-you-die-to-save-itself>
38. <https://www.linkedin.com/pulse/emerging-challenge-ai-self-preservation-instincts-ethical-sarumi-mhugf>
39. <https://www.alignmentforum.org/posts/wnzkjSmrgWZaBa2aC/self-preservation-or-instruction-ambiguity-examining-the>
40. <https://www.csiro.au/-/media/D61/Reports/Artificial-Intelligence-ethics-framework.pdf>
41. <https://www.lesswrong.com/posts/79BPxvSsjzBkiSyTq/agi-safety-and-alignment-at-google-deepmind-a-summary-of>
42. <https://deepmind.google/discover/blog/taking-a-responsible-path-to-agi/>
43. <https://www.industry.gov.au/publications/australias-artificial-intelligence-ethics-principles/australias-ai-ethics-principles>
44. https://en.wikipedia.org/wiki/AL_alignment
45. <https://www.un.org/en/delegate/two-new-mechanisms-promote-cooperation-ai-governance>
46. <https://www.ibm.com/think/topics/ai-alignment>
47. <https://arxiv.org/html/2501.03151v1>
48. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11897388/>
49. <https://www.bcs.org/media/j3rjfncp/white-paper-ai-ethics-and-governance-for-organisational-agility.pdf>
50. <https://millennium-project.org/wp-content/uploads/2023/08/AGI-Governance-Phase-1-1.pdf>
51. https://www.reddit.com/r/singularity/comments/13modal/on_ubi_vs_basic_post_scarcity/
52. <https://openexo.com/insight/ai-driven-abundance-navigating-the-practical-realities-of-a-post-scarcity-era>
53. <https://witness.ai/blog/ai-alignment/>
54. <https://openai.com/safety/how-we-think-about-safety-alignment/>
55. <https://www.alignmentforum.org/posts/EgRJtwQurNzz8CEfJ/dodging-systematic-human-errors-in-scalable-oversight>
56. https://www.reddit.com/r/singularity/comments/1e63stg/stuart_russell_says_ais_will_develop_a/
57. <https://www.anthropic.com/news/claudes-constitution>
58. <https://papers.ssrn.com/sol3/Delivery.cfm/5389563.pdf?abstractid=5389563&mirid=1>

59. <https://www.alignmentforum.org/posts/DqaoPNqhQhwBFqWue/principles-for-picking-practical-interpretability-projects>
60. <https://www.alignmentforum.org/posts/mGCcZnr4WjGjqzX5s/the-checklist-what-succeeding-at-ai-safety-will-involve>